

Basic Dense Matrix Multiplication in Python, Java and C: Effect of Matrix Size on Runtime and Memory

Big Data – Individual Assignment 1 · Grado en Ciencia e Ingeniería de Datos, ULPGC

Amai Suárez Navarro

October 2026

Abstract

We study how the matrix size n affects the runtime and peak memory of the same i-j-k triple-loop dense matrix multiplication implemented in Python, Java and C. All three programs read identical float64 inputs, produce bitwise-identical results and are timed with the same kernel-only boundary on one laptop (Intel Core i5-5350U, 2 cores, 8 GB RAM, macOS). Within each size regime runtime grows close to the cubic model (fitted exponents 3.0–3.3 for $n \leq 192$), but the cost per inner iteration rises about 16-fold for C (9-fold for Java) between $n = 16$ and $n = 1536$, so a single fit over $n \geq 128$ gives exponents near 4.0. Java is 1.3–1.8 $\times$ slower than C for small matrices and equal to C (ratio 1.0) for $n \geq 1024$, where both appear to be limited by memory access. CPython is 110–145 $\times$ slower than C while the data fit in cache and 52–73 $\times$ slower for $n = 256$ –512. Under a budget of 60 s per repetition, the largest feasible size was $n = 512$ for Python and $n = 1536$ for Java and C. These conclusions apply only to the tested implementations, settings and machine.

1. Introduction and objective

Dense matrix multiplication is a basic kernel of data-science workloads and has a simple $O(n^3)$ algorithm, which makes it a convenient case for studying how language runtimes and the memory hierarchy affect performance.

Research question. How does the matrix size n affect the runtime and memory use of the same basic triple-loop multiplication implemented in Python, Java and C, under the documented conditions on one machine?

Hypotheses (stated before running the study).

- **H1** – Runtime grows approximately as n^3 for all three implementations.
- **H2** – C and Java (JIT-compiled to native code) are within a small factor of each other, while CPython is one to two orders of magnitude slower.
- **H3** – The time per inner iteration (t/n^3) increases once the matrices no longer fit in cache, because the inner loop reads B column-wise with a stride of $8n$ bytes.
- **H4** – Peak memory grows as n^2 , with a larger cost per element for Python (boxed floats) and a larger fixed offset for Java (JVM).

2. Background

The classical algorithm computes each of the n^2 output entries as a dot product of length n , i.e. n^3 multiply-add pairs or $2n^3$ floating-point operations. On real hardware the runtime also depends on the memory access pattern. In the i-j-k order, A is read along a row (unit stride) but B is read down a column (stride n elements), which has poor spatial locality and causes cache and TLB misses once B no longer fits in the caches [1], [2]. High-performance libraries avoid this with blocking and packing [3]; these optimisations are deliberately **not** used here and are the subject of Assignment 2.

CPython executes bytecode in an interpreter and represents each float as a separate heap object, so the cost per arithmetic operation is large. Java bytecode is compiled at run time by the HotSpot JIT, so warm-up runs must be separated from steady-state measurements [4]. C is compiled ahead of time.

The rounding error of a length- n dot product in floating point is bounded by $\gamma_n |a|^T |b|$, with $\gamma_n = nu/(1-nu)$ and $u = 2^{-53}$ [5]. This bound is used for the correctness tolerance.

3. Methodology

3.1 Implementations

All three programs implement exactly the same loop nest on row-major storage:

```
for i in 0..n-1:
  for j in 0..n-1:
    s = 0.0
    for k in 0..n-1: s += A[i*n+k] * B[k*n+j]
  C[i*n+j] = s
```

	Python	Java	C
Storage	flat list of Python float objects (pointers to boxed doubles)	contiguous double[]	contiguous double* (malloc)
Toolchain	CPython 3.13.5 (Anaconda)	OpenJDK 21.0.6 (JetBrains Runtime), javac --release 17	Apple clang 14.0.0 (invoked as gcc)
Settings	no NumPy in the kernel; 1 warm-up	-XX:+UseSerialGC -Xms256m -Xmx4g; 3 warm-ups	-O2 -std=c11 -ffp-contract=off; 1 warm-up

The unavoidable difference is storage: the Python list stores pointers to separate float objects instead of contiguous doubles. `-ffp-contract=off` prevents fused multiply-add in C, so the three programs perform the same IEEE-754 operations in the same order. Java ≥ 17 uses strict IEEE-754 semantics and does not fuse operations automatically.

3.2 Inputs

The inputs are square float64 matrices with entries drawn uniformly from $[-1, 1)$. They are generated once by `scripts/gen_matrices.py` using Python's `random.Random` (Mersenne Twister) with the string seed `20262027:n:A` or `20262027:n:B`. They are stored as raw little-endian binary files that all three languages read, so the logical inputs are identical. SHA-256 hashes of every file are recorded in `inputs/manifest.csv`.

Declared sizes: $n = 16, 32, 64, 128, 192, 256, 384, 512, 768, 1024, 1536, 2048$. The list combines powers of two with intermediate values so that both trends and power-of-two effects can be seen.

3.3 Environment

Item	Value
Machine	MacBook Air (Intel), run on 2026-10-08
CPU	Intel(R) Core(TM) i5-5350U CPU @ 1.80GHz (Broadwell-U, 1.8 GHz base, 2.9 GHz turbo)
Cores	2 physical cores, 4 logical processors
Caches [6]	L1d 32 KB and L2 256 KB per core, L3 3 MB shared
RAM	8 GiB
OS	macOS, Darwin kernel 21.6.0 (x86_64)
Python	CPython 3.13.5 (Anaconda); NumPy 2.1.3 for validation only
Java	OpenJDK 21.0.6 (JetBrains Runtime), options <code>-XX:+UseSerialGC -Xms256m -Xmx4g</code>
C compiler	Apple clang version 14.0.0 (clang-1400.0.29.202), flags <code>-O2 -std=c11 -ffp-contract=off</code>

All kernels are single-threaded. The JVM also runs internal threads (JIT compiler, serial GC, housekeeping); CPython runs a single thread. The laptop was plugged in and no other applications were in use. CPU affinity cannot be set on macOS, and the frequency governor is not exposed, so Turbo Boost was active and is not controlled.

3.4 Correctness

`tests/validate.py` runs two groups of checks:

1. **Exact cases** with small integers, which are exactly representable, so the results must match exactly: a hand-computed 2×2 product, 1×1 , $A \cdot I = A$, $I \cdot B = B$, $A \cdot 0 = 0$ and random integer matrices up to $n = 17$.
2. **Random float64 matrices** ($n = 2 \dots 128$) compared against NumPy 2.1.3 as a trusted reference. A result is accepted when $|C - C_{\text{ref}}| \leq 2\gamma_n(|A||B|)$ componentwise.

All 63 checks passed (`results/validation_report.txt`). The largest observed error was $7.1 \cdot 10^{-15}$ at $n = 128$, which is 0.7 % of the bound. The three languages produced bitwise-identical outputs in every random case. During the benchmark, each program printed a checksum (the sequential sum of C) after every repetition. These checksums were identical across languages and repetitions for every tested n , which confirms that the computation was actually executed.

3.5 Measurement protocol

- **Timed region.** Only the kernel (the three loops) is timed. Input loading, allocation of C, the checksum and output are outside the timed region. C needs no zero-initialisation because every entry is written exactly once.
- **Timers.** Monotonic nanosecond timers: `time.perf_counter_ns`, `System.nanoTime` and `clock_gettime(CLOCK_MONOTONIC)`.
- **Batches.** For small n , one sample times $\max(1, \text{round}(2 \cdot 10^7 / n^3))$ consecutive multiplications (for example 4883 at $n = 16$ and 1 for $n \geq 256$). The batch is the same for all languages, and reported times are per multiplication.
- **Repetitions.** One process per (language, n) runs W warm-ups (Python 1, Java 3, C 1) followed by 5 measured repetitions. Warm-ups are kept in the raw CSV and excluded from the statistics. The order of languages is randomised for each size with a fixed seed, to limit bias from drift.
- **Statistics.** Median and interquartile range (IQR) over the 5 repetitions. The empirical exponent b is the least-squares slope of $\log t = a + b \log n$.
- **Memory.** Peak resident set size (RSS) of the whole process, including the interpreter or JVM. On macOS this is `ru_maxrss` of the child process, obtained with `os.wait4`. For Java, the used heap after the run is also recorded. The theoretical storage of the three arrays is $24n^2$ bytes.

- **Budgets (declared before running).** 60 s per repetition, enforced as a process limit of $(W + 5) \cdot 60 \text{ s} + 30 \text{ s}$, and 4 GB of memory. After a timeout, larger sizes are not attempted for that language.

4. Results

All raw data are in `results/run_20261008-202121/`. Tables and figures are generated by `python3 scripts/analyze.py results/run_20261008-202121`.

Table 1. Time per multiplication (median of 5 repetitions), relative spread, cost per inner iteration (median / n^3), throughput, ratio to C and peak RSS.

Method	n	Median	IQR/median	ns/iter	GFLOP/s	$\times C$	Peak RSS (MiB)
Python	16	304.8 μ s	4.4 %	74.41	0.027	109.3	10.2
Python	32	3.69 ms	1.5 %	112.56	0.018	131.7	10.3
Python	64	31.69 ms	3.9 %	120.88	0.017	143.4	10.7
Python	128	260.10 ms	1.2 %	124.02	0.016	133.5	12.1
Python	192	1.04 s	2.9 %	147.63	0.014	143.5	14.4
Python	256	2.30 s	1.5 %	137.02	0.015	64.5	17.7
Python	384	8.66 s	0.8 %	152.89	0.013	73.1	27.1
Python	512	22.92 s	1.2 %	170.74	0.012	51.6	40.3
Java	16	5.1 μ s	0.6 %	1.24	1.610	1.8	36.6
Java	32	39.0 μ s	0.4 %	1.19	1.682	1.4	36.6
Java	64	300.0 μ s	1.0 %	1.14	1.748	1.4	36.9
Java	128	2.45 ms	0.8 %	1.17	1.711	1.3	37.3
Java	192	9.65 ms	0.6 %	1.36	1.467	1.3	38.4
Java	256	37.12 ms	0.6 %	2.21	0.904	1.0	39.7
Java	384	134.79 ms	4.8 %	2.38	0.840	1.1	43.6
Java	512	771.36 ms	10.1 %	5.75	0.348	1.7	49.2
Java	768	1.83 s	1.4 %	4.04	0.495	1.3	63.9
Java	1024	10.43 s	1.6 %	9.72	0.206	1.0	85.5
Java	1536	39.90 s	1.1 %	11.01	0.182	1.0	150.1
C	16	2.8 μ s	5.3 %	0.68	2.938	1.0	0.7
C	32	28.0 μ s	2.6 %	0.85	2.341	1.0	0.7
C	64	220.9 μ s	4.6 %	0.84	2.373	1.0	0.8
C	128	1.95 ms	9.9 %	0.93	2.152	1.0	1.0
C	192	7.28 ms	3.4 %	1.03	1.943	1.0	1.5
C	256	35.63 ms	1.2 %	2.12	0.942	1.0	2.2
C	384	118.47 ms	12.7 %	2.09	0.956	1.0	4.1
C	512	444.32 ms	7.5 %	3.31	0.604	1.0	6.7
C	768	1.37 s	0.6 %	3.02	0.662	1.0	14.2
C	1024	10.81 s	4.4 %	10.07	0.199	1.0	24.7
C	1536	39.72 s	0.7 %	10.96	0.182	1.0	54.7

Table 2. Empirical growth exponent b from a log-log least-squares fit, over the predefined range $n \geq 128$ and piecewise over three size regimes.

Method	$n \geq 128$	$n \text{ 16-192}$	$n \text{ 256-768}$	$n \text{ 1024-1536}$
Python	3.20	3.23	3.31 ($n \leq 512$)	-
Java	3.95	3.02	3.71	3.31
C	4.02	3.14	3.40	3.21

Table 3. Largest feasible size under the declared budgets and the failures observed.


Method	Largest feasible n	Largest attempted n	Failures
Python	512	768	$n=768$: timeout (limit 390 s)
Java	1536	2048	$n=2048$: timeout (limit 510 s)
C	1536	2048	$n=2048$: timeout (limit 390 s)

Table 3b. Repetitions completed before the process limit was reached. They exceed the declared 60 s per-repetition budget and are excluded from all statistics.


Method	n	warm-ups done	measured reps done	time per completed rep [s]
C	2048	1	2	104.2, 104.1
Java	2048	3	2	99.3, 98.2
Python	768	1	3	79.3, 79.9, 78.4

Table 4. Effect of JIT warm-up in Java: first and last warm-up run divided by the median of the measured repetitions.


Method	n	first warm-up / median	last warm-up / median
Java	16	1.80	1.00
Java	32	1.82	1.01
Java	64	1.95	1.00
Java	128	1.63	0.99
Java	192	1.38	0.92
Java	256	1.44	1.03
Java	384	1.17	0.96
Java	512	0.84	1.11
Java	768	0.99	1.02
Java	1024	1.24	1.00
Java	1536	1.00	1.00




Runtime versus n (log-log). Points: median; bars: IQR; dashed line: n^3 reference anchored at the smallest C point.



Median time divided by n^3 (ns per inner-loop iteration). A flat line would mean purely cubic growth.



Peak resident set size of each process versus n , with the theoretical array storage $24n^2$ bytes.



Median runtime of Python and Java relative to C.

5. Discussion

Growth with n (H1, partly supported). Inside each regime the growth is close to cubic: for $n \leq 192$ the fitted exponents are 3.02 (Java), 3.14 (C) and 3.23 (Python). Over the predefined range $n \geq 128$, however, C and Java give $b \approx 4.0$. This does not mean the algorithm changed. The arithmetic work is exactly n^3 , but the cost of each inner iteration is not constant (Fig. 2): for C it rises from 0.68–1.03 ns ($n \leq 192$) to about 2.1 ns ($n = 256$ –384), 3.0–3.3 ns ($n = 512$ –768) and 10–11 ns ($n = 1024$ –1536). The cubic model therefore describes the operation count, while the constant in front of n^3 depends on where the data live in the memory hierarchy.

Memory hierarchy (H3, supported, mechanism partly hypothetical). The steps in Fig. 2 match the column-wise access to B. The step between $n = 192$ and $n = 256$ occurs where B ($8n^2$ bytes: 288 KiB $\rightarrow$ 512 KiB) outgrows the 256 KB per-core L2 cache [6]. The largest step, at $n = 1024$, occurs where even B alone (8 MiB) is far larger than the 3 MB L3 cache, so most accesses go to DRAM. Two observations suggest more than simple capacity effects:

- $n = 512$ is more expensive per iteration than the larger $n = 768$, in both C (3.31 vs 3.02 ns) and Java (5.75 vs 4.04 ns). With $n = 512$ the column stride is exactly 4096 bytes, and power-of-two strides are known to map many accesses onto the same cache sets, which causes conflict misses [2].
- At $n \geq 512$ every step down a column touches a new 4 KiB page, so TLB misses are also likely.

Without hardware performance counters we cannot separate these causes. They are plausible explanations consistent with the data, not measured facts.

Java versus C (H2, supported). After warm-up, Java is 1.3–1.8 $\times$ slower than C for $n \leq 192$ and 1.0–1.3 $\times$ of C for $n = 256$ –768 (1.7 $\times$ at $n = 512$). Java takes the same time as C (ratio 1.0) at $n = 1024$ and 1536, where both reach about 0.18–0.21 GFLOP/s. The convergence is consistent with a memory-bound kernel: once most time is spent waiting for B, differences in generated code matter little. The small- n gap may come from differences in the generated machine code, such as array bounds checks; we did not inspect the JIT output. Table 4 shows why warm-up matters: the first Java warm-up at small n took 1.4–1.95 $\times$ longer than the steady state, while the third was within a few percent of it. The Java result at $n = 512$ is an outlier (ratio 1.7, IQR/median 10 %, warm-ups getting slower). We do not have evidence to explain it.

Python (H2, supported). CPython costs 74–171 ns per inner iteration, against about 1 ns for C in cache. This makes it 110–145 $\times$ slower than C for $n \leq 192$ and 52–73 $\times$ slower for $n = 256$ –512. The ratio falls because C slows down sharply once it becomes memory-bound, while Python’s per-iteration cost is dominated by bytecode interpretation and object handling and grows much less (2.3 $\times$ between $n = 16$ and 512, against 4.9 $\times$ for C over the same range). Python’s own exponent over $n = 128$ –512 is 3.20.

Memory (H4, supported). C’s peak RSS is almost exactly the array storage: 54.7 MiB at $n = 1536$ against 54.0 MiB for three float64 arrays. Java adds a fixed offset of about 36 MiB for the JVM, and at $n = 1536$ it uses about 17 bytes per matrix element. The extra bytes are plausibly the input buffers kept until garbage collection, since used heap was 79 MB at the end of the run. Python starts at about 10 MiB and needs about 40 bytes per element at $n = 512$ (40.3 MiB

in total). This is consistent with an 8-byte list pointer plus a 24-byte float object per element, plus temporary buffers used while loading. Memory was never the limiting factor: the largest process used 241 MB, far below the 4 GB budget.

Limits. The per-repetition budget, not memory, determined the largest feasible size. Python at $n = 768$ needed about 79 s per repetition and C and Java at $n = 2048$ about 98–104 s (Table 3). These partial repetitions are reported but not used in any statistic. Following the cubic trend, one Python repetition at $n = 1536$ would take 10 minutes or more; this is an extrapolation, not a measurement.

Variability and validity. IQR/median was below 5 % in 25 of 30 configurations, with a maximum of 12.7 % (C, $n = 384$). No observations were removed. The main threats to validity are:

- a single laptop with Turbo Boost and thermal behaviour not controlled, and no CPU pinning on macOS;
- one process per configuration, so variation between process launches is not captured;
- the memory metric includes the runtime and is not identical across languages;
- the sizes are coarse, so the exact position of each step is not resolved.

The results do not show that one language is faster in general. They describe these three straightforward implementations on this machine.

6. Conclusions

- For the tested i-j-k implementations, runtime follows the n^3 operation count only within each size regime. The cost per inner iteration grows about 16-fold for C and 9-fold for Java between small matrices and $n \geq 1024$, so the effective exponent over $n = 128$ –1536 is close to 4.
- C and Java behave alike: Java is 1.3–1.8× slower while data are in cache and equal to C for large matrices. CPython is roughly 50–145× slower than C, depending on n .
- Memory grows as n^2 in all three, with about 8, 17 and 40 bytes per element for C, Java and Python respectively, and was never the binding constraint.

Open questions are:

- how much of the large- n cost comes from cache conflicts, cache capacity and TLB misses; this could be tested with hardware counters, or with the unit-stride i-k-j order and blocking (Assignment 2);
- how the kernel scales across threads (Assignment 3).

Use of generative AI

Claude (Anthropic) was used to help plan the experimental design, to write the implementations, benchmarking and analysis scripts, and to draft the text of this report from my measurements. I built and ran all programs, the correctness tests and the benchmark on my own computer. All numbers in this report come from those runs, recorded in the repository. I reviewed the code and the text, and I checked that each reference exists and supports the statement where it is cited.

References

- [1] J. L. Hennessy and D. A. Patterson, *Computer Architecture: A Quantitative Approach*, 6th ed. Cambridge, MA, USA: Morgan Kaufmann, 2017.
- [2] U. Drepper, “What every programmer should know about memory,” Red Hat, Inc., Tech. Rep., 2007.
- [3] K. Goto and R. A. van de Geijn, “Anatomy of high-performance matrix multiplication,” *ACM Trans. Math. Softw.*, vol. 34, no. 3, Art. 12, 2008.
- [4] A. Georges, D. Buytaert, and L. Eeckhout, “Statistically rigorous Java performance evaluation,” in *Proc. OOPSLA*, 2007, pp. 57–76.
- [5] N. J. Higham, *Accuracy and Stability of Numerical Algorithms*, 2nd ed. Philadelphia, PA, USA: SIAM, 2002.
- [6] TechPowerUp, “Intel Core i5-5350U,” CPU Database. [Online]. Available: <https://www.techpowerup.com/cpu-specs/core-i5-5350u.c1812>

Reproducibility materials. GitHub repository <https://github.com/amaisuarez/assignment1>. The study was run from commit `ffcf45a`; raw data in `results/run_20261008-202121/`.